

## 1

### Task

Configure the system to use the following **NTP pools**:

- 0.europe.pool.ntp.org
- 1.europe.pool.ntp.org

Next, change the **timezone** of this system to Europe, Bucharest.

### Solution

To add the NTP pools to the system, edit the /etc/systemd/timesyncd.conf as below:

[Time]

NTP=0.europe.pool.ntp.org 1.europe.pool.ntp.org

Restart the **systemd-timesyncd** to reflect the changes:

```
sudo systemctl restart systemd-timesyncd
```

To change timezone of this system to **Europe, Bucharest** run below command:

```
sudo timedatectl set-timezone Europe/Bucharest
```

=====

## 2

### Task

Add a **cron job** for the user called **john**. Don't use the system-wide crontable, but rather add it to the personal crontable of the user called **john**.

Make sure that this cron job runs every **Wednesday at 4AM**. The command it should execute is `find /home/john/ -type d -empty -delete`.

Switch back to the bob user once the task is done.

### Solution

1) Switch to user john

```
sudo su john
```

2) Edit the crontab using `crontab -e` command as shown below to run jobs

every **Wednesday at 4 AM**:

```
0 4 * * 3 find /home/john/ -type d -empty -delete
```

=====

## 3

### Task

There is a network interface on this system which has the IP address **10.5.5.2** associated with it. What is the name of this **network interface**? Create a file in `/opt/interface.txt` and add a single line to it containing the exact name of that interface.

### Solution

To get the name of network interface which has ip address 10.5.5.2, run below command:

```
ip -o -4 addr list | grep '10.5.5.2' | awk '{print $2}' | sudo tee /opt/interface.txt
```

=====

## 4

### Task

An administrator added a new user called `jane` to this system. But a few mistakes were made. Fix the following problems:

1. The administrator wanted to allow `jane` to run `sudo` commands. But instead of adding "`jane`" to the secondary/supplemental "`sudo`" group, the administrator changed the primary group to `sudo`. Fix this by doing the following: Set the **primary/login group** back to the group called `jane`. And add the user to the **secondary/supplemental group** called `sudo`. In essence the primary group for the user called "`jane`" should be "`jane`". And the secondary group should be "`sudo`".
2. Currently, the **home directory** path for the `jane` user is set correctly. But the directory itself is missing. Fix this by creating the `/home/jane/` directory. Make sure that the directory is owned by the `jane` user and `jane` group.
3. The **default shell** for the user called `jane` is set to `/bin/sh`. Change the default shell to `/bin/bash`.
4. Finally, set the password for `jane` to `1234`.

#### Solution

- 1) Change Primary Group of Jane Back to **jane**:  
`sudo usermod -g jane jane`
- 2) Add the user **jane** to the secondary group **sudo**:  
`sudo usermod -aG sudo jane`
- 3) Create the **Home Directory** for Jane and Set Ownership:  
`sudo mkdir -p /home/jane`  
`sudo chown jane:jane /home/jane`
- 4) Change Default Shell to **/bin/bash**:  
`sudo usermod -s /bin/bash jane`
- 5) Set the Password for Jane:  
`echo 'jane:1234' | sudo chpasswd`

=====

## 5

### Task

Perform the following tasks:

1. Set up a port redirection rule that does the following: it **redirects** any **IPv4** packet from the `10.5.5.0/24` CIDR range incoming on port `81` to another machine with the IP address `192.168.5.2` on port `80`. To simplify this task, you are not required to specify input or output network interfaces.
2. Don't forget to add the proper **masquerading rule** so that packets redirected from `10.5.5.0/24` have a way of getting back to that sender, by our machine sitting in the middle and acting as an ad-hoc router between those two networks.
3. Make sure that after you add the rules you make them **persistent** (so that when the machine is rebooted these changes are not lost).

#### Solution

- 1) Set Up Port Redirection Rule:  
`sudo iptables -t nat -A PREROUTING -p tcp -s 10.5.5.0/24 --dport 81 -j DNAT --to-destination 192.168.5.2:80`
- 2) Add the masquerading rule:  
`sudo iptables -t nat -A POSTROUTING -s 10.5.5.0/24 -j MASQUERADE`
- 3) To make iptables rules persistent:

```
sudo apt install iptables-persistent
```

```
=====
```

6

#### Task

In `/home/bob/certs/` directory you will find 4 files. That's because we generated two self-signed TLS certificates for you. **Delete the 2 files containing the private keys.** But **preserve the certificate files.**

At this point you're left with 2 files containing 2 separate certificates. They both use the RSA algorithm. But one is using 2048 bits for its cryptography purposes, while the other is using 4096. **Delete the certificate that is using 2048 bits.**

#### Solution

1. To find out which files are not a certificate, run the below command, and if you get output as `Unable to load certificate`, then it is a key.

```
openssl x509 -in file* -noout -text
```

2. To find out the certificate that is using the 2048 bit, run the below command:

```
openssl x509 -in file* -noout -text | grep "Public-Key"
```

```
=====
```

7

#### Task

There is a file at `/opt/aclfile`. Currently no one has permissions to read, write, or execute this file, not even root. But instead of working with regular permissions, use **ACL** for this task. Add the following to the access control list:

The user called `janet` should be able to **read and write** to `/opt/aclfile`. Just read and write, no execute permission for this **ACL entry**.

#### Solution

- 1) Set read and write permissions for `janet` on `/opt/aclfile`:

```
sudo setfacl -m u:janet:rw /opt/aclfile
```

```
=====
```

8

#### Task

Add **two security limits** to the configuration of this system:

1. The user called `janet` should have a `hard` limit so that she can't open more than 100 processes.
2. The group called `mail` should have a `soft` limit so users in the group not be able to create files larger than 8192 kilobytes.

#### Solution

To enforce security limits edit the security limits configuration file `/etc/security/limits.conf` as below:

```
janet hard nproc 100
```

```
@mail soft fsize 8192
```

```
=====
```

## Task

In your home directory you will find a subdirectory called **project**. Navigate to it and then do the following:

1. Add **file1** to the staging area to prepare it for a future commit.
2. Commit this file with the exact following message: **Created first required file**.
3. Now upload your changes to the remote repository already associated with your local repository. Everything is already set up for you. Use the remote repository aliased as **origin** and upload the **master** branch and password to push the code is **C0ntr0lplan3Pa\$\$wd**.

## Solution

1) Move into the **/home/bob/project** directory

2) Run below command to create a file:

```
touch file1
```

3) Add the **file1** to staging by using **git add .** command

4) Commit the changes with the message **Created first required file**:

```
git commit -m "Created first required file"
```

5) Push changes to the **master** branch on remote repo

```
git push origin master
```

=====

## 10

## Task

Perform the following tasks related to **SELinux** on **node01**:

1. First, check if SELinux is running in **enforcing**, **permissive**, or **disabled** mode. Create the file **/opt/selinuxmode.txt**. And write your answer to that file. Just one line where you write a single word: **enforcing**, **permissive**, or **disabled**, according to the status you found.
2. There is a file that has the **wrong SELinux type label**. Please correct that and restore the **default SELinux label** for the file at **/usr/bin/less**.

Credentials to access **node01**:

**Name:** bob

**Password:** caleston123

## Solution

1) Check the current SELinux status on **node01** and store it in file **/opt/selinuxmode.txt**:

```
getenforce | sudo tee /opt/selinuxmode.txt
```

2) Correct SELinux Type Label for **/usr/bin/less**:

```
sudo restorecon -v /usr/bin/less
```

=====

## 11

## Task

**Nginx** is installed but it's **not running** on **caleston-lp10**. Make the necessary changes so that:

1. Nginx is started immediately.
2. And also, Nginx will start up automatically every time the system boots up.

After starting Nginx, it has spawned at least **3** processes. We are not interested in the **master** process, only the **worker** processes. Under what **username** are

these **worker** processes running? Create a new file at `/opt/nginxuser.txt` and add a single line to it with that username, other than **bob**.

#### Solution

1) Start Nginx Immediately:

```
sudo systemctl start nginx
```

2) Enable Nginx to start on boot:

```
sudo systemctl enable nginx
```

3) List the Nginx processes and store the username at `/opt/nginxuser.txt`:

```
ps -ef | grep nginx
```

4) Copy the username of worker processes of nginx and store it in `/opt/nginxuser.txt`

```
=====
```

## 12

#### Task

A basic **LVM** structure exists on the `node01`. Make some changes to it:

The volume group called `volume1` currently only includes `/dev/vdb`.

Add `/dev/vdc` to `volume1`.

We have a logical volume called `lv1`. Resize this logical volume to 2GB`.

Credentials to access `node01`:

**Name:** bob

**Password:** caleston123

#### Solution

1) Add `/dev/vdc` to `volume1`:

```
sudo vgextend volume1 /dev/vdc
```

2) Resize the Logical Volume `lv1` to 2GB:

```
sudo lvresize --size 2G /dev/volume1/lv1
```

```
=====
```

## 13

#### Task

Perform the following tasks on `node01`:

1. Install **Git**.
2. In your home directory, download the entire repository from <https://github.com/htop-dev/htop>. By default, this action should create a new subdirectory called **htop** where all the project's files are located.
3. Switch to the **htop** subdirectory. Follow the project's instructions to **build** (compile) the htop application.
4. Install the newly built htop application. By default, the application should be installed in `/usr/local`, in the `bin` subdirectory

Credentials to access `node01`:

**Name:** bob

**Password:** caleston123

Switch back to `caleston-lp10` once the task is done

#### Solution

1) Install Git on the `node01` server:

```
sudo apt install -y git
```

2) Clone the `htop` repository into your home directory and navigate into the local repo:

```
git clone https://github.com/htop-dev/htop
cd htop
3) Install the dependencies of htop:
sudo apt update
sudo apt install libncursesw5-dev autotools-dev autoconf automake
build-essential
4) Install htop by running the following commands:
./autogen.sh
./configure
make
sudo make install
```

=====

**14**

#### Task

Create a **virtual machine** with the following parameters on caleston-lp10:

1. For the **operating system information parameter** use the ubuntu22.04.
2. Name the virtual machine mockexam2.
3. Assign 1024 MB of RAM to this machine.
4. Assign it one virtual CPU.
5. Import the disk image /var/lib/libvirt/images/ubuntu.img to this virtual machine.
6. At the end of your command, you can add the parameter --noautoconsole to avoid waiting for this virtual machine to boot up, and not get stuck in the virtual console after it initializes.

After you create this virtual machine, run a separate command to make mockexam2 automatically start up every time the system boots up.

#### Solution

Run the below command to launch vm with ubuntu22.04 image:

```
virt-install \
--name mockexam2 \
--ram 1024 \
--vcpus=1 \
--os-variant=ubuntu22.04 \
--import \
--disk path=/var/lib/libvirt/images/ubuntu.img \
--noautoconsole
```

Configure the virtual machine to start automatically on boot:

```
virsh autostart mockexam2
```

=====

**15**

#### Task

Create a network bridge between these network interfaces: [ "eth1", and "eth2" ] on node01. Call the bridge bridge1. Turn **DHCP off** for IPv4 for both interfaces. However, for the **bridge** itself, turn **DHCP** on for **IPv4**.

Credentials to access node01:

**Name:** bob

**Password:** caleston123

Solution

To create a bridge network with eth1 and eth2, create a file `/etc/netplan/99-bridge.yaml` and edit the file as below:

network:

```
version: 2
renderer: networkd
ethernets:
  eth1:
    dhcp4: no
  eth2:
    dhcp4: no
bridges:
  bridge1:
    dhcp4: yes
    interfaces:
      - eth1
      - eth2
```

Apply the new network configuration using the `netplan apply` command

=====

**16**

Task

Perform the following tasks on node01:

1. Remove the Docker image called `nginx`.
2. Start a container based on the `httpd` image. **Name** it `apache_container`. Instruct Docker to redirect connections coming to port 80 on the host to port 80 of this container. Also instruct Docker to **restart** this container **only on failure**, with a **maximum number of retries** set to 3 (you will have to look through the correct manual to find the parameter you need).

Credentials to access node01:

**Name:** bob

**Password:** caleston123

Switch back to `caleston-lp10` once the task is done

Solution

1) Delete the containers using `nginx` image:

```
docker rm nginx_container
```

2) Delete the `nginx` image:

```
docker rmi nginx
```

3) Run the below command to create `httpd` container:

```
docker run --detach --publish 80:80 --name=apache_container --
restart=on-failure:3 httpd
```

=====

**17**

Task

The node01 was configured to use LDAP entries from a certain server. However, some configuration options are wrong. Edit the correct configuration files and fix the following mistakes:

1. Our name service local daemon is configured to look for an **LDAP server** at the **wrong IP address** (currently **10.9.9.8**). Fix this and configure the correct IP which is: **192.168.121.167**. Make sure your changes become active after you edit the configuration file.
2. Our system is currently configured to get **group data** and **password data** from the LDAP server but not user data. Configure it to get **user data** as well.

Credentials to access node01:

**Name:** bob

**Password:** caleston123

Switch back to caleston-lp10 once the task is done

Solution

1. Edit /etc/nslcd.conf to change the ldap server:

```
# The user and group nslcd should run as.
uid nslcd
gid nslcd
```

```
# The location at which the LDAP server(s) should be reachable.
uri ldap://192.168.121.167/ #updated
```

```
# The search base that will be used for all queries.
base dc=example,dc=org
```

```
# The LDAP protocol version to use.
#ldap_version 3
```

```
# The DN to bind with for normal lookups.
#binddn cn=anonymous,dc=example,dc=net
#bindpw secret
```

2. Edit /etc/nsswitch.conf to get user data:

```
# /etc/nsswitch.conf
#
# Example configuration of GNU Name Service Switch functionality.
# If you have the `glibc-doc-reference' and `info' packages
installed, try:
# `info libc "Name Service Switch"' for information about this file.
```

```
passwd:      files systemd ldap #updated
group:       files systemd ldap
shadow:      files ldap
gshadow:     files
```

3. Restart the nslcd service:

```
sudo systemctl restart nslcd
```





